

# Chromadb vs Pinecone

## CHROMADB (Local)

- ✓ Runs entirely local
- ✓ Free and open source
- ✓ No API costs
- ✓ Fast for small datasets
- ✗ Less scalable
- ✗ No multi-user support

## PINECONE (Cloud)

- ✓ Managed service
- ✓ Massive scale
- ✓ Multi-tenancy built-in
- ✗ Needs internet
- ✗ Costs \$0.18 per 100K vectors
- ✗ Data leaves infrastructure

## DECISION:

For MVP: Chromadb as the backend for RAG pipeline

Re-evaluate in phase 2

Impact: Save ~50-200/month

Timeline: 3 days vs 2 days

Risk: Need migration path later

